

Screening and Recruiting Talent At Teacher Colleges

Sebastian Gallegos
IDB

Christopher Neilson
Princeton University and NBER

Franco Calle
Princeton University

November 18, 2019

Latest Version Downloadable: [here](#)

In this paper, we study policies that use pre-college academic achievement to restrict or incentivize entry to teacher-colleges. First, we use historical records of college entrance exam scores since 1967 in Chile and administrative data on the population of teachers to document a robust positive and concave relationship between several measures of teacher productivity and teachers' own pre-college academic achievement. We then study the effectiveness of two policies that used pre-college achievement to recruit or screen out students entering teacher-colleges. Using a regression discontinuity design based on the government's recruitment efforts, we evaluate the effectiveness of targeted scholarships at shifting career choices of high achieving students as well as the effect on the overall stock of teachers predicted effectiveness. We then study the effects of a recent screening policy that forces teacher colleges to exclude below-average students. We retroactively simulate this policy rule and evaluate its success at screening out low performing teachers. We compare this benchmark policy rule to a series of potential data-driven policy rules and we find that even simple screening policies can identify a significant portion of ex-post low performing teachers. In both policies studied, screening low performing students is more effective than targeting recruitment efforts to only very high achieving students. These findings suggest that the combination of better data and flexible prediction methods can be used to implement practical screening policies in some contexts.

PRELIMINARY AND INCOMPLETE. PLEASE DO NOT CITE.

*Gallegos: Inter-American Development Bank, 1300 New York Avenue, Washington, DC. Neilson (corresponding author) and Calle: Department of Economics, Princeton University, 20 Washington Road, Princeton, NJ 08544. Email: cneilson@princeton.edu. We thank the DEMRE and Ministry of Education (MINEDUC) of the government of Chile for access to the sources of data and help locating and accessing data records. We thank Macarena Alvarado and Alvaro Carril for contributions to early versions of this paper and seminar participants at The University of Chicago, Princeton's Education Research Section, the Workshop on Education and Child Development, University of Stavanger, Norway, the Education Seminar at the Inter-American Development Bank, and the APPAM and SREE annual meetings. Gallegos also thanks the Center for Studies of Conflict and Social Cohesion (CONICYT/FONDAP/15130009). All errors are our own, and the views expressed in this paper are those of the authors and should not be attributed to the Inter-American Development Bank or any other institution. Online Appendix available at <https://christopherneilson.github.io/work/teacher-quality.html>.

1 Introduction

A common policy objective of governments all over the world is to increase the productivity of the stock of teachers. One set of policies looks to increase the effectiveness of current teachers through incentives, increasing their effort through accountability measures and rewards. For instance, Hoxby and Leigh (2004) documented that talented teachers were leaving the profession due to the compression of wages, and Biasi (2018) finds that a policy based on rewards improved effective teaching. These policies typically require measuring teacher effectiveness carefully and tracking it over time to write optimal incentive compatible contracts making them difficult to rely solely on incentives to promote effective teachers in many contexts, especially in poorer, developing countries. A second complementary set of policies looks to promote the recruitment of talented individuals into the teaching profession. Policies such as college tuition scholarships make attending teacher college more attractive, and screening entry policies can exclude candidates from teaching colleges that do not meet minimal standards. Designing both of these policies efficiently depends crucially on pre-college academic achievement being an important determinant of how effective (or not) a student will be as a teacher in the future. In the US context, there is consensus that predicting talented teachers' ex-ante is notoriously difficult, making it impractical to design systems relying solely on high-powered benefits to recruit talented individuals or minimal standards to exclude low performing candidates. In developing countries, the relevant margin of quality seems to be closer to whether teachers show up for work (Kremer et al., 2005; Muralidharan et al., 2014; Duflo et al., 2015), though some studies suggest that pay for performance schemes could work for those already in the teacher workforce (Mbiti et al., 2019; Muralidharan and Sundararaman, 2011).

In this paper, we study policies that use pre-college academic achievement to 1) screen out applicants that are unlikely to become effective teachers later in their careers and 2) target resources to recruit high achieving students to become teachers. We first show evidence that in the context of a developing country, pre-college academic achievement in school and in college entrance exams is systematically correlated with teacher performance. We use recently digitalized historical data on college entrance exams from 1967 onward merged with administrative records on the population of teachers in Chile. We document a robust positive and concave relationship between teachers' pre-college academic achievement and a variety of short, medium and long run teacher outcome measures. Short and medium run outcomes include the probability of graduation from teacher college, college exit exams, and employment and wages in schools. Long run measures of productivity include government teacher evaluations, student test scores, and school value added as well as own teacher value added. Broadly we find that below average students are systematically associated with lower performance as teachers when their performance is measured up to thirty and forty years later. Our findings suggest that at least in the context of a developing economy such as Chile, own academic readiness can be seen as a necessary condition. College training nor experience seem to be enough to undo the initial deficiencies in the relationship between entrance exam and measures of performance.

We then evaluate two policies that look to affect the distribution pre-college academic achievement of the students entering teaching colleges in the country. The first is a policy implemented to attract high scoring students by offering tuition subsidies at participating teacher colleges. The policy also required teacher colleges to voluntarily exclude students with below average entrance exam scores in order to be eligible for the program. The second policy, which was implemented later, forced all teacher colleges to exclude students with below average college entrance exams with an allowance for students with high ranking GPA. We show that these policies had a significant impact on both individual college major choices of students as well as an effect on the aggregate distribution of predicted teacher quality.

To evaluate the first recruiting policy enacted in 2011, we use a regression discontinuity design to show how high-powered incentives such as scholarships can attract high achieving students to the teaching profession. Choice probabilities for students with scores one standard deviation above the mean jumped 40% at the cutoff, but on a small base leading to an increase of 1,000 (out of 20,000) additional students with high academic achievement. Approximately fifty percent of teacher colleges participated and had to exclude the majority of low performing applicants. This did significantly increase average pre-college academic achievement by cutting off a significant part of the lower tale of the distribution and as a result, raised predicted teacher effectiveness. Higher achieving students graduated and performed as predicted on exit exams six years later lending credence to the idea that the relationships documented above are invariant to this type of policy. However, several years later, tuition was made free for many other programs due to a different policy and the same regression discontinuity showed that the effectiveness of the policy was significantly diminished. This result serves as a reminder that recruitment incentives are only as good as the next best option and that high achieving students have many good alternatives, so it is harder to move them towards teaching.

We then turn to study a second policy that used pre-college academic achievement as a direct screening policy enacted in 2017. It excluded low performing students from entering teaching colleges independently of whether the college wanted to participate. We evaluate this very recent policy by replicating it going backward and identifying who would have been affected and whether they did indeed become ineffective teachers. Partial equilibrium analysis shows that if implemented, these rules would have bound 25% of students entering teaching colleges in 2016 and would have affected 20% of current teachers, including 87% of the worst performers based on government teacher evaluations. We then evaluate the scope to improve the screening by using standard machine learning methods to determine the predictive content of the information available on academic achievement prior to entry to college. We train a standard model that classifies potential teachers based on a rich set of precollege test information. The results suggest policy-makers can use initial characteristics to determine college specific cutoff criteria that excludes the same number of students as any of the screening policies proposed by the government, but a better job could be done of selecting the students that are more likely to be higher performing teachers in the future. Partial equilibrium analysis shows that –when comparing with the most strict screening rule proposed by the government– our data driven rule would have increased the number of students graduating in time in around 6%, increased the number of teachers working after 7 years of being enrolled in college in 10%, and increased the number of teachers working in well performing schools according to SIMCE in 10%.

Taken together, these results indicate that policies that use pre-college academic achievement can be useful to increase the predicted quality of teachers. It is also the case that data driven adjustments to the policies could make them more effective. In addition, given we find that predicting lower performing teachers is easier than high performers, screening policies seem the most viable policy in this context. Screening policies could possibly be complemented with targeted subsidies to students in a broad “competent category” rather than attempting to recruit the very high achieving students.

These results are important because they have direct policy implications. If teacher effectiveness, or lack thereof, is possible to predict early on, policies could focus resources on recruiting and retaining the most promising candidates and filtering out applicants who are more likely to become ineffective teachers. This is particularly important because teacher labor markets are known to be inefficient, mis-allocation of talent can be widespread and in many cases, and there is limited scope to sideline or retrain ineffective teachers once they are in the system, especially in the public sector (see, e.g., Estrada (2019) for the Mexican case).

The results presented in this paper highlight avenues for further research in an increasingly data-rich environment where prediction is a key input to policy design. While entrance exam scores, transcripts and GPA remain an imperfect measure of teacher quality, they are broadly available and can be used to target resources and design policy at zero marginal cost. In many developing countries, this may be the only option they have to design policy, as implementing the infrastructure needed to implement repeated student tests or regular teacher classroom evaluations is not available. Our findings suggest that at least in the context of low to middle income countries such as Chile, resources that look to subsidize teacher training should be targeted towards prospective teachers that have a minimal level of baseline academic achievement. Own academic readiness can be seen as a necessary condition, though perhaps not sufficient.

These results are related to recent work in developing countries that shows that teacher value added is important and varies a lot. For example Bau and Das (2018), show evidence on teacher value added from Pakistan finding important mis-allocation of pay in the public sector, and in contrast to our results, that teacher quality value added is not explained by teacher observable characteristics. Araujo et al. (2016) estimate teacher value added randomly assigning students to teachers and find significant variation in teacher quality, and that observational measures of quality were correlated with estimated value added. Similarly, to Bau and Das (2018), they do not find teacher characteristics to be correlated with teacher value added. The results presented in this paper seem potentially at odds with this evidence and that from the US (Rockoff, 2004; Rothstein, 2006). However, there are important differences worth mentioning is the empirical exercise. In our case, we are using rich pre-college academic achievement for the population of teachers which may not be available to researchers in other contexts. Another important note is that we find most of the correlation comes from very low achieving students who become teachers and this margin may not be relevant for more developed countries.

2 Context, Policy and Data

2.1 Context

Chile is a Latin American country with almost universal levels of educational enrollment in primary (99%) and secondary school (92%), which are well above other countries of the region (OECD, 2009) and similar to developed countries. The country is in the advanced stages of a demographic transition, with low fertility and mortality rates, and relatively high life expectancy (World Bank, 2011). Consequently, enrollment in primary and secondary education has plateaued and even showed a slight decrease over the last ten years (from 3.1 million in 2008 to 2.9 million in 2018). In the meantime, the number of classroom teachers¹ has increased from 125,000 in 2008 to 164,000 in 2018 (MINEDUC, 2019), which has led to a reduced student-teacher ratio (from 25 to 18) that is close to the levels displayed by OECD countries (World Bank, 2013). In addition, teacher absenteeism is estimated at 5% (Paredes et al., 2015) which is much lower than other countries in earlier stages of development (e.g., Chaudhury et al. (2006) estimate an average of 19 % for Bangladesh, Ecuador, India, Indonesia, Peru and Uganda). With enough teachers in the classrooms, the policy focus switched in the last ten years to bring more qualified individuals to the teaching profession.

Attracting more skilled individuals to be teachers is a challenge because, among other factors,

¹Teachers work in the different types of schools which differ in their funding and administration. Public schools are funded and administered by the government; voucher schools are funded mainly with public funds but administered by privates; and private schools are both funded and administered privately.

teachers are paid less (Mizala and Nopo, 2016) compared to other professionals.² Consistently, we know from the related literature that college graduates with higher college entrance scores are less likely to enter teaching (Manski, 1985; Hanushek and Pace, 1995; Vegas et al., 2001), and Chile is no exception. Figure 1 shows that between 2007 and 2010 students enrolled in fields other than education (engineering, law, medicine, etc.) scored about 1 standard deviation (σ) above the national mean in the college entrance exam, while teacher college students scored only 0.1σ above. In addition, we computed that the scores from teacher colleges have been declining over time; in 1995 students from teacher colleges scored 0.3σ over the national mean. This pattern is similar to the evidence for the U.S. (Bacolod, 2006; Corcoran et al., 2004; Podgursky et al., 2004; Hoxby and Leigh, 2004).

2.2 Policies to Recruit Teachers In Chile

The Chilean government implemented in 2011 a scholarship aiming to recruit higher-scoring students to the teaching profession and also to restrict access for the lower-scoring students. The policy, called *Beca Vocacion Profesor (BVP)*, consisted in full tuition subsidies for prospective students who scored about 1σ above the mean in the college entrance exam. The policy also required participating teacher colleges to reject applicants with scores below the national mean. In 2017, the government implemented a screening policy, imposing new requirements for admissions at all teacher colleges across the board. The government policy required applicants to teacher colleges to have college entrance exam scores at least as high as the median of the distribution of test-takers or have a high-school GPA in the top 30% of their high school graduating cohort.

Figure 2 shows the evolution of the average college entrance exam for freshmen in education, and is suggestive of the policy effects on scores in 2011 and 2017. The Figure shows the percentage increase in PSU scores for freshmen in the education, health and STEM fields, from 2007 to 2018, taking 2007 as the base year. There is a sharp increase in the scores for freshmen in education from year 2010 to 2011, and another from 2016 to 2017, which coincide with the implementation of the BVP and the new government rule. At the same time, Figure 2 show no variations in the scores achieved by freshmen in STEM or health fields.

Also, Figure 3a shows how the policy changes flattened the increasing trend of students enrolled in education from 2011 to 2018, meanwhile, for health the increasing trend continued from 2004 until 2018. The trend for enrollment in teacher colleges as freshman correlates with the shift in the number of teaching programs available. Figure 3b shows that after 2011 the number of programs available sharply declined, with a last and very step reduction in 2017. In the case of health, for instance, the number of programs available maintained its increasing pattern with some plateau period between 2013 and 2015 but retaking the increasing trend until 2018.

2.3 Data on Pre-College Academic Achievement

The main measure of teachers' pre-college academic achievement that we use in this paper is their scores on college entrance exams taken since 1967. These data have been collected as a part of the work done in Hastings et al. (2014) where the authors collected digital copies of old books and newspapers, digitalizing test score data back to the first test in 1967.

The Chilean national college entrance exam is similar to the SAT in the U.S. with additional subject tests beyond math and language. Currently the exam is called the *Prueba de Selección*

²Mizala and Nopo (2016) show that teachers are underpaid in Latin America compared to other professionals and technicians, after controlling by observable characteristics linked to productivity. In Chile in particular, the underpayment for teachers was about 18% in 2007.

Universitaria (PSU) and has been administered once a year since 2004. Prior to that a similar test called *Prueba de Aptitud Académica* had been implemented from 2003 back to 1967, which makes Chile to have one of the longest running centralized college assignment systems in the world.³ Test-takers complete exams in mathematics and language as well as other specialized subjects. The scores are scaled to a distribution with a mean and median of 500 and standard deviation of 110. The exam scores are required to apply to all public universities and most private universities and institutes.

2.4 Data on Teacher Productivity

Chile has developed important efforts to measure teacher performance. According to a recent review by World Bank, Chile has the most advanced system of teacher performance evaluation in Latin America (Bruns and Luque, 2015). The most important assessments are exit exams for graduates from teacher colleges, government evaluations for public school teachers and employment at schools measured as good performers by “*Sistema de Medición de la Calidad de la Educación*” (SIMCE).

In addition, the Ministry of Education of the Government of Chile (MINEDUC) keeps administrative records with information on the population of teachers at all schools including private and government subsidized schools since 2003. In spite of the general richness in the data, what is not readily available are measures of teacher value added because students take standardized tests on irregular grades. To fill this gap, we use value added measures collected for a sample of teachers in De Gregorio et al. (2019), where students took tests at the beginning and at the end of the academic year to capture value added of the teacher by controlling for initial conditions. As part of the same project students are also surveyed on their perceptions regarding whether their teacher is effective. We use all these sources of information to measure teacher quality and describe them briefly below.

Our measures of teacher productivity span over earlier outcomes (at age 23) to longer run outcomes measured more than thirty years later. In particular, our teacher outcomes include short run outcomes such as graduation from teacher colleges and college exit exams; and long run outcomes such as earnings, employment, external classroom teaching evaluations, and student’s value added and perception. We describe briefly each dataset below. In our online appendix we describe each measure in detail.

2.4.1 Administrative Data Sources

Graduation from Teacher Colleges. We use information from enrollment in teacher colleges for years 2004 to 2009 for about 85K individuals, which we link to graduation records from years 2009 to 2017. We study on time graduation rates (within 5 years after initial enrollment, at 23 years old) and late graduation (up to 8 years after enrollment, at 26 years old).

Exit Exams. Our data consists in microdata from all the *INICIA* and *Diagnostica* exam test-takers between 2009 and 2017. The sample consist of about 35K just graduated teachers with scores in a disciplinary knowledge test (e.g., math knowledge for math teachers) and a pedagogical knowledge test (e.g., capacity of explaining concepts in a coherent way). At the time of the exam test-takers were approximately 25 years old on average.

Government Evaluations. We gathered information for 63K classroom teachers in public schools, evaluated between 2004 (when the ED started) to 2017. On average, teachers were 40 years old at

³A detailed explanation on the application and enrollment process for the period 1980-2009 is presented in Hastings et al. (2014) and a review comparing centralized systems in the world in Neilson (2019)

the time of the evaluation. They have on average 12.5 years of tenure (years working in schools).

Employment in Schools. We gathered information for about 240K graduates from teacher colleges in years 1995 to 2017 and merged with the population of employed teachers between 2003 to 2018. We compute whether graduates work ever as teachers and we also study whether they work as teachers 2, 5, 10 years, and 10, 15 and 20 years later, respectively and correlate that with entrance exam scores. The age at employment after ten years and twenty of graduation average 37 and 46 years old respectively.

Wages in Schools. Teachers with information on wages are 37 years old, 70 percent female, with 9 years of tenure. Teachers working in public schools are 38 percent of the sample. They benefit from a special labor code, which makes wages grow with tenure and not expected to change with productivity. However, the voucher sector operates under the regular and more flexible labor code, and thus teacher wages can be given a market clearing interpretation, associated to productivity. They represent 62 percent of our sample.

The data on employment and wages comes from administrative records on the population of teachers, which is formally reported by the school principal to the MINEDUC at the beginning of each academic year. The data, available for the period 2004-2017, contain 3.4 million records which correspond to about 370,000 unique teachers. This dataset has information on teachers' demographics (gender, age), specialization level (preschool, primary, secondary) and subject taught (math, language, general, other), as well as tenure (years working as a teacher) and job characteristics (hours of contract, type of contract and starting in 2011, wages).

2.4.2 Collected Data Sources

Value Added. Despite its rich administrative records, Chile does not implement tests designed to measure value added. To overcome this lack of data we implemented those tests for a sample of 9th, 10th and 11th graders and their teachers. The tests measured math skills at the beginning and the end of the academic year, and were specially designed to measure the gain in student achievement during that period.⁴

Student Perceptions. As part of the same project described above, we collected information from students' perceptions regarding effective teaching. The perceptions survey contains 38 questions and follows the recommendations of the Measures of Effective Teaching study carried out in the U.S. (Kane and Cantrell, 2010). Questions are categorized eight dimensions of teaching practices and classroom environment: positive culture and learning environment, student understanding checked for and ensured, engaging learning environment, expectations held by teacher, student input and ideas valued, learning fully internalized by students, encouraging and supporting relationships fostered, and classroom participation. We use students' perceptions in a as an additional measure of teacher quality.

3 Pre-college Academic Achievement and Measures of Teacher Productivity

In this section we document the systematic correlation between pre-college academic ability most of the teacher productivity measures described in the previous section. We estimate parametric

⁴The tests were developed by the Measurement Center of the Catholic University of Chile (MIDE UC) and are called SEPA tests. The name SEPA comes from the acronym in Spanish of Learning Progress Assessment System (*Sistema de Evaluación de Progreso del Aprendizaje*). In De Gregorio et al. (2019) we use these learning gains to study the effectiveness of traditional teachers vis-a-vis "Ensen Chile" teachers. "Ensen Chile" teachers are recent college graduates recruited to teach for two years in low-income communities, resembling the U.S. experience of "Teach for America" teachers.

regressions of teacher outcomes at different moments of their careers on their own entrance exam scores taken at age 18. We also describe the empirical relationship showing non-parametric plots leveraging on our large sample sizes.

The general takeaway is that the empirical relationship between pre-college skills as a student and teacher productivity later on is positive and concave. In Table 1 we show the coefficients of 14 separate regressions for different measures of teacher performance on the college entrance exam score (in standard deviation units) and its square. The coefficients on scores are positive and significant, and most coefficients on the square are negative.

In Figure 4 and Figure 5 we examine early outcomes of students from teacher colleges, like exit exams and graduation rates. We find that college entrance exams scores are positively correlated with the exit examinations college students of pedagogy take before graduating. According to Table 1, one standard deviation on the test scores that current teachers took years ago, is associated to an increase of 0.50σ on both the disciplinary and pedagogical skills in the exit exams respectively; and an increase in 0.46σ and 1.27σ in writing skills and ICT skills (available for smaller samples). The same pattern can be visualized in Figure 4 for the pedagogical and disciplinary tests. Non-parametric plots for ICT and writing tests are shown in the online appendix.

With regards graduation, the relationship appears concave for both graduation after 5 and 8 years after enrollment. The results show that an increase in one standard deviation on the college entrance exam scores leads to an increase in graduation rates after 5 (8) years of enrollment in teachers colleges of 7.3 (11.8) percentage points (relative to a baseline graduation rate of 34.7% (47.3%). One hypothesis that might explain the concave relation is that exceptional students might either switch to another career and quit the teaching profession to a more attractive career with potential higher expected income, or they might just have an attractive outside option in the labor market.

The next set of results are for later outcomes, when individuals are teaching and working in schools. Figure 6 show the bivariate relation between college entry exams scores and teacher evaluations taken up to 30 years later. The relationship is concave again, suggesting that early scores may have a higher potential for identifying low performance teachers than high performing ones thirty years later. Coefficients in Table 1 show that an increase of one standard deviation in entry exam scores translate into an increase of 0.62σ and 0.48σ on the teacher evaluation score overall and portfolio score respectively.

Table 1 is consistent with the concave productivity story we were telling in the Figures, and corroborates a non linear pattern between scores in PSU and the probability of working as a teacher years after graduation. First, an increase of one SD in psu scores increases the likelihood of working as a teacher in 38% percentage points (5 years after graduation) relative to a baseline of 44%. Nevertheless, a significant fraction of teachers in the right tail of the distribution of college preparedness quit the profession by that time.

Figure 8 shows how hourly wages vary with scores, by teachers working in public and private schools. The slope is much steeper for teachers working in the private sector, and rather flat for teachers working in the public sector. Consistently, the coefficients in Table 1 show that a standard deviation increase in scores is associated to 0.26σ and 0.45σ of hourly wages for teachers working in the private and public sector respectively. The magnitude of the coefficient over wages is more prominent for the sample of teachers in the private sector since private schools can move salaries unrestrictedly as teacher productivity changes, the same dynamic does not occur in the public sector where wages are less flexible and determined primarily by the years of service in the public sector which is not a concise measure of productivity.

In the next sections of the paper we study two policies that look to screen out low performing students from teachers colleges or to attract high achieving students to teachers colleges. The first

policy called *Beca Vocacion Profesor* is a policy implemented in 2010 which was aimed primarily at recruiting high scoring students into teacher colleges by offering full scholarships, stipends and even paid semesters abroad. Short run effects show significant effects on choices made by students as many enrolled in teacher colleges instead of alternative options. Seven years later, the medium run results on early indicators of teacher productivity strongly suggest that the relationship between pre-college academic achievement and teacher productivity is invariant to the policy studied. The second policy was a mandatory screening policy implemented in 2017 and requires accredited teachers' colleges to accept and allow enrollment only to students with entrance exam scores above the mean. We study the predicted change in teacher productivity given the policy and simulate the partial equilibrium effects of the policy had it been implemented in the past.

4 A Carrot and Sticks Approach to Recruiting and Screening

This section presents results of the *Beca Vocacion Profesor* (BVP) policy, which we briefly described in subsection 2.2. The results of this "carrots and sticks" policy was that the proportion of high achieving students rose by approximately 50%, although enrollment from the lower end of the distribution continued high at non participating institutions. We provide details on specifics of the policy next, and then show results on college participation choices and students outcomes.

4.1 BVP Policy Specifics

The BVP was first implemented in the academic year 2011,⁵ This policy gave full scholarships and other incentives such as stipends and paid semesters abroad for students with scores from approximately the highest 30% of the admissions test distribution, who matriculate at teaching colleges. Specific subsidies for students were determined exclusively by the average math and language college entrance exam score (PSU test). Students who scored above $\mu + \sigma$ (600 points, top 25%) were eligible for a full tuition scholarship. If they scored above $\mu + 2\sigma$ (700 points, top 5%), they were eligible for a full tuition scholarship and a monthly stipend of approx \$US150 (approx 50% minimum wage). If the student scored above $\mu + 2.2\sigma$ (720 points, top 2%), in addition to a full tuition scholarship and a monthly stipend, they were eligible for a paid semester abroad at a prestigious teaching college. Advertisements mentioned a semester abroad at Stanford or in Finland to name a few.

In addition, the policy imposed colleges to screen out low performing students. In particular, teacher colleges were required to implement a cutoff score at the 50th percentile of the average score distribution if they wanted their top students to matriculate with the BVP scholarship.⁷ Requirements for colleges included being an accredited program for at least 2 years at all campuses. This is determined by the National Commission of Accreditation (CNA). The teaching college must also maintain a minimum score of 500 (the average score) with no more than 15% exceptions among the entering class starting in 2011.

⁵Requirements for students included having applied and been admitted to an eligible teaching college as a new first year student in 2011 with an entrance exam score from December of 2010. Students previously enrolled in teaching careers are not eligible for this particular modality of the scholarship.

⁶If the student had obtained another scholarship called *Beca Excelencia Academica* the cutoff will be 580. These are very few and we ignore students with this scholarship in the analysis.

⁷The colleges were allowed to accepting up to 15% of their matriculation below that cutoff

4.2 Descriptive Effects

On Aggregated College Participation Choices. College participation is a key feature of this policy and its effectiveness. The policy offers scholarships to students which is beneficial to colleges competing for students. It also allows colleges to add eligibility to their marketing strategy which can be a signal of quality given it is costly to participate without having a high ability student body. The result was that participation was high among public institutions (forced in fact) and traditional private colleges but participation was very low among institutions with less qualified students and Professional Institutes that also offer teaching degrees. Eligible institutions and careers covered only 40% of matriculated students in 2010. Approximately 1/3 career/college combinations that were eligible did not participate and approximately 1/4 career/college combinations were not even eligible. In general it is safe to assume the screening aspect of the policy had less impact than it could have been due to low participation of teachers colleges, where only 50% decided to join the program. Figure 9 shows the estimate of participation probability for eligible teacher colleges conditional on the average scores of their freshman class the year prior. Public institutions were mandated to participate leading to a positive probability even for very low scores by the figure clearly shows participation being driven by the chosen policy cutoffs.

On Aggregated Students Choices. In the application process to the 2011 BVP, the number of applicants to the teacher scholarship was 28,179. Among these, only 3,385 (12%) were eligible for the benefit and 3,252 accepted it (96% of the ones that received it).

Among the ones that accepted the fellowship, 3,063 are beneficiaries of the type I fellowship⁸, which is the focus of this paper. The effect the policy had on the distribution of scores is quite striking. Figure 10 shows the distribution of test scores for students matriculated in teaching careers and non teaching careers in 2009-2010 and compares the same distributions with 2011 when the policy was in place. Notice the large shift in the density after 600 points which is the first cutoff. Also notice the drop in density below 500 which is due to the restriction on colleges participating (that do not admit students below this threshold). The left panel of Figure 10 shows how the proportion of incoming teacher college students changed at specific intervals of interest.

4.3 Causal Effects on Individual Students Choices

Leveraging the large amount of data, we first plot the conditional choice probabilities and then turn to a formal analysis using a regression discontinuity design. Using a regression discontinuity empirical design we show that the probability of choosing to matriculate in a teaching college rose by 40% to 200% at different point of the distribution. The effects were found to be greater for students who were male and who came from public schools or families with lower income levels.

We estimate the probability as a proportion of students that chose a teaching career between a group of students with similar PSU score. We present in Appendix A the same estimation conditional on the type of school. If we analyze all the teaching careers in the Educational System, we can see a decrease in the probability of studying a teaching career under 500 points and also an increase above 600 points, reaching the maximum level in that score, to then decrease when the PSU score increases.

If we analyze only the probability of choosing an eligible teaching career, the probability for 2011 decreases under the average μ (500 point average of math and language) almost to zero.

⁸The government implemented two types of fellowships, one for the students that took the standardized test the year before and the other for students that were pursuing careers related to teaching careers and wanted to be teachers. Less than 5% of the winners of the fellowship were applicants through the second type.

To approximate the impact of the fellowship using only the temporal variation, we calculate the change in the probability of choosing a teaching career in several parts of the distribution. Figure 12 shows that there exists an increase close to 40% in the probability of enrollment to a teaching career around $\mu + \sigma$ (600 points), which increases above 100% at $\mu + 2\sigma$ (700 points). Under μ (500 points) we observe a decrease of 25%. The effect is much larger under the 500 points if we consider only the eligible teaching careers as can be seen in the right panel of Figure 11. In this case, the probability of studying these careers increase in 40-50% when 600 points is reached, then increase to more than 200% when 700 points is reached. Under 500 points we can observe the probability decrease in almost 100%, which is explain by the restriction of cutoff of 500 points for the careers.

To explore the hypothesis of heterogeneous effects of the fellowship, we observe how the enrollment probability for a teaching career changes for students that came from different types of schools.

$$Pr(E_d | PSU_i = 600, X_i) = Pr(V(i, E_d) > V(i, j) \forall j_{available} | SchoolType_i, PSU_i)$$

We found a similar pattern in each case. However, the levels are different because the student that came from voucher schools has a higher probability of choosing a teaching career than the students from a private school. Moreover, this difference is even bigger when we compare it with public schools. In terms of the effects of the tuition scholarship, we observe that in all the cases the probability increases if the student has more than 600 points and decreases if the student gets less than 500 points.

4.4 Regression Discontinuity Design

The program in their eligibility criterion uses a discontinuity that can be used to evaluate potential causal effects over different outcomes. The eligibility criterion is a decision rule that fix a cutoff point to select beneficiaries and non-beneficiaries for the program. This analysis allows estimating the impact of the fellowship around a cutoff point. This is more informative than the non-parametric analysis, but still has some problems when interpreting the results because the analyzed outcome is a continuous non-parametric function. We will study the effect of the fellowship over the extensive margin by studying if the program attracted and therefore increased the number of teachers in the labor market with scores above the threshold of 600 points, and also if the financial incentives operated through the intensive margin by increasing the likelihood –conditional on scores– of graduating and entering early into the labor force as teachers for students with scores above 600, 700 and 720 points. This methodology does not use many assumptions, but is valid only locally around the threshold points.

The intuition behind this method, although it is possible that exists a lot of unobserved heterogeneity in the way the student takes their decision, is the fact that students with similar scores are on average equals. This is, the cutoff point could be approximated to a random neighborhood, in other words, the students that are just above the cutoff point are similar to the student that are just below the cutoff point which requires that both populations are close enough of the cutoff point to secure comparability. The intuition is that students above and below these cutoffs are on average identical so that the noise associated with the test score process effectively give us a local experiment.

Graphically we observe the effect around the 600 points cutoff in Figure 13. The points are the average among students in groups of 5 points of difference and the line to each side of 600 is a linear regression. Figure 13b shows that the overall enrollment into universities didn't change for students with scores above the threshold; meanwhile Figure 13a shows that the threshold crossing

effect increased the likelihood of enrolling into a teaching program from 12% to 18%. The combination of those results suggest that there was a shift in the margin of choosing a career rather than in the number of students that were qualified to apply for the scholarship in the margin. Students that would have preferred a non teaching career in the margin, ended enrolling into teaching colleges because of the financial incentive.

Although, we could not find any significant effect of the third level of benefits over the probability of enrollment in a teaching career we found important effects over the less generous incentives. Particularly, as Table 2 shows, the increase in the probability in 600 is of 5.3 percentage points on a base of 12, which implies an increase on the percentage close to 35%. Around 700 points we observe an impact of 2.7 percentage points on a base of 2.8. This is equivalent to duplicate the probability of choosing a teaching career. If we consider only the eligible careers, the result is the same, which explains why of this margin, most of the enrollees are in the eligible careers.

4.4.1 *Persistent Effects*

Seven years later it can be seen that this policy has had lasting effects. Early productivity indicators such as graduation rates, exit exams and employment probabilities suggest that the policy boosted positive outcomes and raised the quality of students who entered into the teaching profession. In addition, these medium run effects validate the prediction models given that the relationship between pre-college academic Achievement and medium run outcomes is invariant to the recruiting and screening policies.

RDD Estimates on Later Outcomes: Dropout, Graduation, Employment, Exit Exams

Moreover, as Figure 14 shows, there is an important contribution of the BVP policy on the extensive margin by marginally increasing the number of students with scores above the threshold in 20 students on average or 34%, in 24 the number of students working in schools after 7 years of enrollment or 52%, and in 20 the number of students working in good schools for those or 74%. However, it seems that the financial incentive did not have an actual impact over the likelihood of succeeding in any of the outcomes mentioned before, conditional scores, as can be seen in all of the figures in the right panel. This means that even when the policy helped recruit more qualified students in the margin, the recruitment program turned being policy invariant because it did not have an effect on accelerating or improving the labor outcomes for those who received the fellowship.

Policy effects depend crucially on other policy

The results mentioned before depend however on the existence of other financial restrictions outside teaching programs. In particular, in 2016 many colleges in Chile became free of tuition and this would potentially reduce the financial incentive generated by the *Beca Vocacion Profesor* policy over enrolling into teaching programs. When examining the policy discontinuity over time, the threshold crossing effect over recruitment drops consistently until 2016 when gratuidad was implemented. Figure 15 plots the percentage increase in the freshmen enrollment rate in teaching colleges attributable to the threshold crossing effect until 2017. The exercise shows that the effect decreased considerably over time from around 35% in 2011 to almost zero percent in 2017.

4.5 Discussion of BVP Policy

At the aggregate level, the policy increases the average quality of predicted teacher effectiveness.

Overall, it is shown that the majority of the effects on the average student entering teaching college were due to the incentives to screen out low performing students and not the recruitment effects on their own. The reason is that while free tuition and additional subsidies raised choice probabilities significantly, the number of students at those margins is small relative to the population of students in teaching colleges. We estimate that a total of 1000 additional students entered teaching colleges from the top 30% of the distribution but the total number of freshman students at teachers' colleges is closer to 20000 and screening restrictions are estimated to have reduced the bottom tail of the distribution by 4000. In terms of predicted productivity, the concavity of relation between pre-college academic achievement and predicted teacher effectiveness suggests that reducing low performers may be more effective than recruiting very high performers because of the uncertainty on the right tail of the distribution.

5 A Mandatory Screening Policy

This section studies a mandatory screening policy implemented in 2017. The Chilean Government enacted the *NUEVA LEY DE CARRERA DOCENTE* (NLCD)⁹. This is a broad policy that creates a new system of professional development for teachers in the country. The Law includes specific guidelines for the recruitment, development and retention of teachers. One important aspect of the policy is that only accredited programs can now confer teaching degrees. To be accredited, all teacher colleges are required to implement an entrance exam and the exit exam within the last twelve months before students graduate. To be accredited colleges must also comply with new admissions screening rules based on pre-college academic achievement, specifically making use of college entrance exams and high school GPA rank within school.

In this section we explore the consequences of the screening policies in two distinct ways. The first is to apply the screening policy to prior cohorts and determine what the partial equilibrium effects would have been in terms of total matriculation and also to categorize the students that were screened out or in by their ex-post outcomes. In this exercise we expect the screening rule to be successful at blocking entry to teaching colleges to students who did not graduate, did not get teaching jobs, or went on to become less effective teachers. We would expect a successful screening rule to also not leave out students who went on to be highly effective teachers as well.

The second is to train and search among a family of flexible models a data driven method to screen out potentially low performing teachers by relying only on pre-college human capital characteristics. We evaluate how successful our data driven method can be –compared to the screening criteria proposed by the government– at minimizing the mistakes of rejecting future effective teachers while maintaining fixed the number of ruled out students. Our hypothesis is that our data driven method may be able to explode non linearities existent in the data that could be useful for identifying bad performers that would have been ignored by a less flexible rule.

5.1 NLCD Policy Specifics

The requirements for the screening policy affects admissions to all teacher colleges and are designed to be implemented gradually. During the first three years (2017-2019), the screening policy (P17) requires students to either have achieved an entrance exam score above the 50th percentile of the distribution when averaging math and language. Alternatively, students can also avoid the screening rule if their high school GPA is above the 70th percentile within their high school

⁹The law 20.903 is available here [Law 20,903](#)

in their graduating cohort. For the admissions cycle of 2020 to 2022 (P20), the screening rule increases the requirements. Students must have achieved an entrance exam score above the 525 points when averaging math and language scores or have GPA above the 80th percentile. In addition, if students have a GPA above the 60th percentile and test scores in math and language that average above 500 points, they may also matriculate in teacher colleges. Finally, in 2023 and onwards (P23), the screening policy reaches its steady state and requires students to have entrance exam scores above 550 points (the 70th percentile) or be in the top 10% of their cohort GPA. If the student is in the top 30% of the GPA distribution at their high school **and** has scores above the average, then that student can also matriculate at teaching colleges.

All of these conditions are designed as minimal requirements for admission to teacher colleges. Each institution is allowed to consider stricter conditions, define number of vacancies or slots and application mechanisms. However, all the requirements must be informed before the beginning of admission process each year.

5.2 Aggregate Effects on Teacher College Matriculation

As shown by Figure 16, the screening policy that was implemented in 2017 contributed to a significant reduction of the number of students that were below the cutoff of 500 points, however some of them were still eligible since their relative high school ranking was compliant with the requirements of the government.

In Figure 17a we simulated the policy rule of 2017 backwards from 2007 to 2016 to evaluate the number of students that would have been accepted by the policy versus the share of students that applied but were rejected for not complying with the requirements. The trend of the policy performance shows that if the policy were implemented, the number of accepted students would have continue reducing as well as the share of students that would have been rejected for not meeting the policy requirements. According to the figure, the share of students with low ranking and low PSU were between 30 and 31% in 2007 and remained in this range until 2010. From 2011 onwards the trend continued to decrease significantly due to the introduction of BVP and the continuous incorporation of teaching schools into the BVP scheme year by year. The trend, however, shifted in 2016 after the introduction of gratuidad that possibly removed incentives for teaching schools to comply with the requirements asked by BVP.

In Figure 17b, we studied what happened with selected labor outcomes of past teaching students classified as rejected or accepted by the 2017 policy. As can be seen in the Figure 17b, students that would have been rejected by the 2017 policy performed lower in all measures of labor outcomes proposed. For instance, only 10% of the students that classified as rejected by the 2017 policy were likely to be classified as satisfactory in the Exit Exam, 74% lower than the probability for the average accepted student; 29% graduated on time (within 6 years after matriculation), 10% less than the average student accepted; moreover, only 24% of the rejected started working as teachers after 7 years and only 64% of them worked, meanwhile the average accepted student were 38% likely to be working after 7 years and 75% of them in good schools; finally, only 12% of the rejected students were classified as good teachers by the portfolio examination, half as likely as the accepted teachers.

5.3 Simulating the Screening Rule Back In Time

5.4 Towards A More Efficient Screening Policy

Increased availability of data and algorithms can help policy-makers make better predictions. In this case, a better prediction model can help guide the design of a more efficient data driven

screening policy. Building on evidence presented above, and comparing the performance of the screening policies proposed by the MINEDUC we now try to improve on the simulated outcomes produced by the current policy.

We designed our model to be policy invariant by restricting the model to use input features that do not depend on outcomes of the policy; therefore we sacrificed college level characteristics that were strong predictors of the outcome variable but were affected by enrollment results of the policy if implemented in the future. In particular, our predictive features are scores on the entrance exam (all of them) the relative ranking of students in their last year of high school, the average and the standard deviation of the scores in language and mathematics for the teacher's high school as well as the average level of education of the parents for the students evaluated by SIMCE and the average income of the families. Our main outcome variable is being hired by a high school classified as good performer according to SIMCE in less than 7 years of being enrolled in a teaching college.

To highlight the potential role for a flexible screening rule, Figure 18 shows the conditional mean of being hired by a satisfactory school as classified by SIMCE conditional on math and language entrance exam scores. The simple average of math and language scores would cut across this space in a linear way but it can be seen that level curves are quite nonlinear with areas of low probability within the upper left and lower right corners. This suggests a modification of the rule that puts equal weights on math and language at an arbitrary cutoff is likely going to have less success than a more flexible rule.

To estimate and validate our data driven screening strategy, we selected a sample of teachers that allowed us to maximize the number of observations while keeping the maximum number of variables that had a meaningful contribution to the prediction success. In specific, we kept the sample of students that enrolled into a teaching program from 2007 to 2011 primarily because the teachers' registry is only available until 2018 meaning we could only match the dependent variable with students enrolled until 2011, 7 years before they start working. Additionally, we used this sample because we could not match observations for the years before 2007 with high school level variables of SIMCE. Finally, we dropped all observations without individual math, verbal and NEM scores available in the sample. The resulting sample consists of 36990 teachers enrolled in teaching programs, which was split in a training set (80%) and test set (20%) consisting of 29592 and 7398 observations respectively.

After implementing crossvalidation a grid search of models and hyperparameters over the training, we selected the Light Gradient Boosting Machine which reported an area under the curve of 67%, higher than the standard levels of AUC for predicting behavioral outcomes Chalfin et al. (2016).

In Figure Figure 19a we computed how much does our machine learning screening rule improve the labor outcomes of students that would have been accepted by the government screening policies. For evaluating the contribution of the ML, we imposed a probability cutoff for the ML policy that rejects the same number of students that each of the government policies would have rejected in the whole sample. Results of the exercise suggest that the ML screening rule outperforms the government rules by increasing each year the number of students graduating from teacher colleges before 6 years from enrollment in 2% for P17 and around 4% for P20 and P23 (75 and 150 students per year, respectively). The ML rule would have also increased the number of teachers working in schools 7 years after enrollment in around 4% (174 more students each year) for P17 and between 6% and 8% (260 or 347 more students each year) for P20 and P23; finally, the ML rule would also increment in 7% (228 more students) the number of P17 screened in students

working in a good school, and 8% and 10% (260 or 325 more students yearly, respectively) for the cases of P20 and P23.

Even after removing socioeconomic variables at the high school level and relying only on pre college human capital characteristics the model does a good job. Figure 19b shows that even when the data driven method does not outperform the P17 rule by much, it significantly improves the outcomes of P23 which is more strict by definition. The data driven method increases the number of students graduating in time in 6%, and in 10% both the number of students working at schools and in good schools after 7 years.

Although our model outperforms all government’s screening rules, it has some limitations. In Figure 20, we examine up to what extent is it possible to continue increasing the accuracy of our model. For this, we plotted the area under the curve evaluated in the test sample for models trained with different levels of sample size in the training sample. Figure 20a shows that there is still room for improvement in terms of accuracy if the model includes more data to train parameters; however, the contribution seems to have a concave pattern and could be reaching a limit soon. Also, we plotted the contribution of each variable over performance of the model, Figure 20b shows that the most important variable for predicting teaching performance may be the average score for students in high school, however it seems to contribute less than all exit exam scores altogether.

5.5 Discussion of Empirical Results

In the previous sections we have found evidence that students with low levels of pre-college academic achievement tend to be less effective teachers but high scoring students, while likely to be above average teachers, are similar in expectation to above average students. These correlations imply that pre-college academic achievement seems to be a necessary but not sufficient condition to be a great teacher and a not having a minimal level of preparation is predictive of being a less effective teacher later on.

The recruiting and screening policies presented suggest that incentives matter in bringing talent to the teacher profession but if targeted at very high scoring students, the aggregate effects are minimal due to strong outside options. When new government policies expanded free college to other majors, the incentives of the recruiting policy became ineffective.

Overall the main driving force moving the average teacher predicted quality seems to have been screening policies that limit entry into the teaching profession. We show that these policies can be made more effective if the rules governing them are data driven and more flexible.

6 Conclusions

Data are becoming more abundant as administrative sources become available, historical information gets digitalized and new information gets recorded in more detail than ever before. Combined with the development of improved algorithms, these data are lowering the cost of making increasingly accurate predictions and are influencing decisions, such as hiring, in many markets (Agrawal et al., 2018). In this paper we put together historical datasets, administrative records and data sources collected in the field to show that better data and flexible prediction methods can be used to implement enhanced teacher screening policies.

In particular, we first show evidence that in the context of a developing country, pre-college academic achievement is systematically related to a series of measures of long run teacher productivity. We then evaluate the effectiveness of two recent policies that looked to screen or recruit

based on pre-college academic performance and find that both raise the predicted quality of teachers, own graduation rates and exit exams scores. We argue the results indicate that policies that use pre-college academic achievement to either screen out or incentivize an application to teacher colleges can be feasible and useful in some contexts. This is particularly important because teacher labor markets are known to be inefficient in most countries (Neal, 2011) and, in many cases, there is limited scope to sideline or retrain ineffective teachers once they are in the system. If teacher effectiveness was possible to predict early on, policies could focus resources on recruiting and retaining the most promising candidates and filtering out applicants who are more likely to be ineffective teachers later on.

In our analysis we show that while identifying very effective teachers remains challenging, very simple models can use the data to categorize and identify a significant proportion of the worst performing teachers. We find a concave relationship between pre-college academic achievement and later teacher productivity, which we interpret as evidence that in a developing country context such as Chile, basic academic competency is a necessary condition to be an effective teacher. We provide suggestive evidence that this relationship between pre-college academic achievement and productivity is seemingly not caused by high scoring students having differential access to more selective and more effective teaching colleges. In fact, we find no meaningful differences across different teaching colleges on exit exam scores once conditioning on pre-college academic achievement. Taken together, our evidence seems to suggest that neither training or experience are enough to undo the significant initial deficiencies that very low performing students have and thus systematically more likely to be low performing teachers thirty and forty years later when observed in the classroom.

We then evaluate two policies implemented in Chile that look to shape the pool of students entering teaching colleges by screening out low performing students or setting incentives for high performing students based on their pre-college academic achievement.

The first policy, implemented in 2011, aimed at bringing more qualified students into teacher training programs via financial incentives. At the same time, teacher colleges willing participate in the policy had to limit admission to only students with college entrance exams above average. We show evidence that financial incentives such as scholarships and stipends do indeed lead to higher recruitment of talented students entering participating teacher colleges. The probability that a student with 1σ above the mean would choose to enroll at teacher colleges rose by 40% , over a relatively low base of 10%. Moreover, early productivity indicators measured seven years later, show that those talented students have indeed higher graduation rates, exit exams and employment probabilities, as predicted by their higher college entrance exam scores. This piece of evidence suggest that the relationship between pre-college academic ability and later outcomes is invariant to these types of policies and lends credence to policies using college entrance exam scores as predictors of future performance.

We also show that about half of the teacher colleges decided to participate, which reduced significantly the amount of low performing students matriculating in teacher colleges nationwide. We estimate that screening restrictions decreased the bottom tail of the distribution by one fifth of the total freshmen enrollment (4,000 over 20,000 students). In addition, many higher education options became tuition free as part of another government policy years later (2016). This new policy changed relative prices and generated suggestive evidence helping to disentangle effects attributed to the components described above. In practice, we find that the effectiveness of the financial incentives was reduced significantly. The results suggest that inducing colleges to voluntarily exclude the lowest performing students was the most effective aspect of the policy. The results also highlight that the effectiveness of targeting highly talented students with recruiting efforts is highly context dependent and expensive because they have many other valuable options.

A second screening policy implemented in 2017 barred all teaching colleges from admitting students with below average scores unless they had very high GPA. This policy was found to be effective at shifting the predicted quality of future teachers, screening out between 20-30% of the least academically prepared applicants. To evaluate the policy relevance of a minimum standard for entering teaching colleges, we develop a model that classifies potential teacher productivity based on the rich set of pre-college information including GPA course transcripts and entrance exam scores. This model provides feasible cutoff rules that exclude students with a higher chance of being a low performing teacher. Partial equilibrium analysis shows that if implemented, these rules would have been 10% more successful than the screening method proposed by the government by using only pre college human capital characteristics. We interpret these results as suggestive that screening policies can be improved with even simple models and a data driven policy rule.

In both policies studied, the most effective aspect of the policy comes with screening policies aimed at excluding prospective students with scores below the median and not with recruiting the highest ability students. This is both a function of the higher ability to identify low productivity teachers from the bottom of the academic achievement distribution and that recruiting high ability students is hard. Taken together this suggests that increasing the predicted productivity of a cohort of future teachers can be increased first by excluding the lower tail of distribution of academic achievement and potentially using any resources saved to incentivize a large group of simply above average students to enroll in teaching colleges with the former being the more effective of the two.

The evidence presented in this paper could be viewed as being at odds with the consensus in the US literature that teacher effectiveness is not very predictable Rothstein (2015). Part of the discrepancy could be due to better data availability in Chile at this time. In our view, while this is likely, the more relevant issue is that the context is very different and the distribution of underlying academic preparation in a country like Chile is shifted to the left and has a wider dispersion relative to the USA. Exit exams in Chile indicate that more than 40% of teachers do not pass a subject test on their own material they are supposed to teach. It is quite possible that in a context such as the US, the relationship between pre-college academic achievement and later effectiveness is already on the flat part of Figure 6 and these results are mostly relevant for developing countries.

The policy relevance of screening policies for developing countries can be seen as limited because they are usually trying to expand educational access and need more teachers and more classrooms. The temptation to expand rapidly is to lower standards and fill classrooms with bodies. However, recruiting newly minted teachers with potentially low ability and limited prospects will seemingly continue to be as such for the next thirty or forty years as we found in the data for Chile. In this context it might do a country well to consider growing more slowly, sticking with minimal standards for entry into the teaching profession and higher wages, with a mind to planning ahead towards a smoother transition from a system that provides quantity to one that provides quality. Teaching is a relatively cheap degree to offer and supply expanded faster than any other option in Chile after government backed loans were provided by the government for the first time. Many students with low scores then find themselves with limited options but teaching is virtually always feasible for them. Minimal standards for entry or for access to subsidies, can also help regulate the supply of degrees that are being oversupplied by reducing demand from groups that less likely to benefit from those studies.

In this paper we have outlined that screening and recruiting policies implemented before candidates enter college could be feasible and useful in some context. A data driven approach to determining the specific details of the policies seem promising. Future work should study the equilibrium effects of these policies as they will likely affect incentives for universities.

References

- Agrawal, Ajay, Joshua Gans, and Avi Goldfarb**, *Prediction machines: the simple economics of artificial intelligence*, Harvard Business Press, 2018.
- Araujo, M Caridad, Pedro Carneiro, Yyannu Cruz-Aguayo, and Norbert Schady**, "Teacher quality and learning outcomes in kindergarten," *The Quarterly Journal of Economics*, 2016, 131 (3), 1415–1453.
- Bacolod, Marigee**, "Do Alternative Opportunities Matter? The Role of Female Labor Markets in the Decline of Teacher Quality," Working Papers, U.S. Census Bureau, Center for Economic Studies 2006.
- Bau, Natalie and Jishnu Das**, "Teacher Value-Added in a Low-Income Country," *American Economic Journal: Economic Policy*, 2018.
- Biasi, Barbara**, "The Labor Market for Teachers Under Different Pay Schemes," NBER Working Papers 24813, National Bureau of Economic Research, Inc July 2018.
- Bruns, Barbara and Javier Luque**, *Great Teachers: How to Raise Student Learning in Latin America and the Caribbean*, Washington, DC: World Bank, 2015.
- Chalfin, Aaron, Oren Danieli, Andrew Hillis, Zubin Jelveh, Michael Luca, Jens Ludwig, and Sendhil Mullainathan**, "Productivity and selection of human capital with machine learning," *American Economic Review*, 2016, 106 (5), 124–27.
- Chaudhury, Nazmul, Jeffrey Hammer, Michael Kremer, Karthik Muralidharan, and F. Halsey Rogers**, "Missing in Action: Teacher and Health Worker Absence in Developing Countries," *Journal of Economic Perspectives*, Winter 2006, 20 (1), 91–116.
- Corcoran, Sean, William Evans, and Robert M. Schwab**, "Changing Labor-Market Opportunities for Women and the Quality of Teachers, 1957-2000," *American Economic Review*, 2004, 94 (2), 230–235.
- Duflo, Esther, Pascaline Dupas, and Michael Kremer**, "School governance, teacher incentives, and pupil–teacher ratios: Experimental evidence from Kenyan primary schools," *Journal of Public Economics*, 2015, 123 (C), 92–110.
- Estrada, Ricardo**, "Rules versus Discretion in Public Service: Teacher Hiring in Mexico," *Journal of Labor Economics*, 2019, 37 (2), 545–579.
- Gregorio, Soledad De, Sebastian Gallegos, and Christopher Neilson**, "Decomposing Teachers Contributions to Learning," Working Paper, Mimeo 2019.
- Hanushek, Eric and Richard R. Pace**, "Who chooses to teach (and why)?," *Economics of Education Review*, 1995, 14 (2), 101–117.
- Hastings, Justine, Christopher Neilson, and Seth Zimmerman**, "Are Some Degrees Worth More than Others? Evidence from college admission cutoffs in Chile," Working Paper 19241, NBER 2014.
- Hoxby, Caroline and Andrew Leigh**, "Pulled Away or Pushed Out? Explaining the Decline of Teacher Aptitude in the United States," *American Economic Review*, 2004, 94 (2), 236–240.

- Kane, T. and S. Cantrell**, “Learning about teaching research report: Initial findings from the measures of effective teaching project,” Technical Report 2010.
- Kremer, Michael, Karthik Muralidharan, Nazmul Chaudhury, Jeffrey Hammers, and F Halsey Rogers**, “Teacher Absence in India: A Snapshot,” *Journal of the European Economic Association*, Vol. 3, pp 658-67, 2005.
- Manski, Charles**, “Academic Ability, Earnings, and the Decision to Become a Teacher: Evidence From the National Longitudinal Study of the High School Class of 1972,” NBER Working Papers 1539, National Bureau of Economic Research, Inc 1985.
- Mbiti, Isaac, Mauricio Romero, and Youdi Schipper**, “Designing Effective Teacher Performance Pay Programs: Experimental Evidence from Tanzania,” NBER Working Papers 25903, National Bureau of Economic Research, Inc May 2019.
- MINEDUC**, “Sistema de Informacion General de Estudiantes: Idoneidad Docente,” Data File, Retrieved from <http://datos.mineduc.cl/dashboards/19732/bases-de-datos-de-cargos-docentes/>, Santiago, Chile: Ministerio de Educacion de Chile 2019.
- Mizala, Alejandra and Hugo Nopo**, “Measuring the relative pay of school teachers in Latin America 1997–2007,” *International Journal of Educational Development*, 2016, 47 (C), 20–32.
- Muralidharan, Karthik and Venkatesh Sundararaman**, “Teacher Performance Pay: Experimental Evidence from India,” *Journal of Political Economy*, 2011, 119 (1), 39–77.
- , **Jishnu Das, Alaka Holla, and Aakash Mohpal**, “The Fiscal Cost of Weak Governance: Evidence from Teacher Absence in India,” *NBER working paper 20299*, 2014.
- Neal, Derek**, “Chapter 6 - The Design of Performance Pay in Education,” in Eric A. Hanushek, Stephen Machin, and Ludger Woessmann, eds., *Handbook of The Economics of Education*, Vol. 4 of *Handbook of the Economics of Education*, Elsevier, 2011, pp. 495 – 550.
- Neilson, Christopher A.**, “The Rise of Centralized Mechanisms in Education Markets Around the World,” Technical Report August 2019.
- OECD**, *Education at a Glance 2009: OECD Indicators*, Paris: Organisation for Economic Co-operation and Development, 2009.
- Paredes, Valentina, Andreas Aron, and Alvaro Carril**, “Where is the Teacher? Short-run effect of teacher Absenteeism on Student Achievement,” Technical Report, Department of Economics, Universidad de Chile. Mimeo. 2015.
- Podgursky, Michael, Ryan Monroe, and Donald Watson**, “The academic quality of public school teachers: an analysis of entry and exit behavior,” *Economics of Education Review*, 2004, 23 (5), 507–518.
- Rockoff, Jonah E.**, “The Impact of Individual Teachers on Student Achievement: Evidence from Panel Data,” *American Economic Review*, 2004, 94 (2), 247–252.
- Rothstein, Jesse**, “Teacher Quality Policy When Supply Matters,” *American Economic Review*, 105(1): 100-130, 2015.

Rothstein, Jesse M., “Good Principals or Good Peers? Parental Valuation of School Characteristics, Tiebout Equilibrium, and the Incentive Effects of Competition among Jurisdictions,” *American Economic Review*, September 2006, 96 (4), 1333–1350.

Vegas, Emiliana, Richard Murnane, and John Willet, “From High School to Teaching: Many Steps, Who Makes It?,” *Teachers College Record*, 2001, 103 (3), 427–449.

World Bank, *Population aging: Is Latin America Ready?*, Washington, DC: David Cotlear (editor), 2011.

—, “System Approach for Better Education Results (SABER): What Matters Most in Teacher Policies? A Framework for Building a More Effective Teaching Profession. Washington, DC,” Technical Report 2013.